

Using Machine Learning to tell if a Spider Died

Jack Bosco

February 14, 2023

Goal

- Want to learn about circadian rhythms in spiders
 - How does a usual (12L:12D) cycle compare to a (12D:12D) cycle?
- We have large sets of data to help us determine this

```
1111100000001111110001101010100011111011000010010011010101000101111011010100000
1110001111000101000111111001101010001001111001100001001111101000011100000011
00111001110000001001111101111010100100111000110110101101010010100001001010000
111111110111110001100000010011101001111101100111100110101110011111000111001
000101010011100001111100001000011011111100001110000100000111110001101111
000011010011101011101111001001100001010010001110001001000100001001000100001011001
110101110000000101000010100101010011010110101111011100100000001100010110000
0001100100110010111101111011110000101011000010101010100011001010001101010100
10001000000010110100101010010011011001010100111011001011110000110010001000011
00111001101100111100000101010010001000101000011010101001010101110010111010101
000110101101010000101011110110011000111011011010000111100000000011000001101100
000100100111111101000011001010000100101010001100110000110001001100011010111111
011110101000111100100110000010110101010010001101111000000100110000000011110001
00110111000010111100001101010100101111101111010000110101010010001101010101110
00000100010011100010110011000001011010001101101101000100000111111001100101000
010101110110010101010100000011010111010000010011010100011101001000110110010010
0010101000010101001110001110011000111110100001111000011100110010001101001001001
00111000010111000001101011110100100001100101110010000100110010011100001101111111
10111110000111000101000001010010011110010001010110101011110100110100001101101
000111100000101101111110010001011110010011000110010011011101000011100010010
10000001000011110000100110001101001101010110111101101111010010110100101101101
000011110101001110101110011100101100000011011111001101101110100001111110110011
0000100111110011110000100001110111001001100100001100100011001010001011001100
1011010111101010001000000110011111001000001011010011001110011011000010101000111
01100111100101010101111001110010001001100100010011101010110011110100101001010
001001001010100100011011100111011000101111001101100100110000000111010110000100
01010100000000010111101101000110110011011001011010101110100001010000000101100
111100110010101000100011100010101111100100010110001101010001011001010000100000
10011100100100001000011011110110011001001100011000100011010111101000100001100
0001011010001010101011110011000110101010010001101001110000101011000010010011111
01100000010110001101101100100101100101100110001010110111011011001010101
```

Problem

Null Character

Inaction (spider not doing anything) is coded similarly to sleeping.

Actually, its exactly the same

- As is, no human can look at this data and determine if a spider was affected by change in day/night cycle or if it died overnight
- This biases the data, rendering tons of information inconclusive

Problem Redefined

Dead|Not-Dead

A classification algorithm *should* be able to use pattern matching and see when or if a spider died.

For an input string of n 0s or 1s...

- Action Space: n . We should be able to point (somewhat accurately) to the moment when the spider died.
- State Space: 2^n . Every possible arrangement of 0s and 1s needs to be classifiable. This number is HUGE.

Our state and action space makes training **Dead|Not-Dead** non-trivial.

Can we do it?



Yes

If the problem is solvable, then a perceptron can learn to solve it




Not So Fast...

Just because it can doesn't mean it won't take forever

So not only do we need to train an algorithm that solves this problem, we need to be smart about it as to not end up spending \$1,345,232,923 on AWS.

Solution 1: 'Meh' Classifier

Multilayered Perception (MLP)



- a sigmoidal (non-linear) activation function
- outputs on a gradient
 - $0.0 \rightarrow$ The spider died (earlier)
 - $1.0 \rightarrow$ The the spider is totally still alive
 - $0.25 < x < 0.75 \rightarrow$ 'Meh'

Solution 1: 'Meh' Classifier

Pros

- Simple, understandable, we've had these since 1989
- I can (and have) implemented this with only numpy
- quick/inexpensive

Cons

- results may be uninformative
- Vanishing Gradient Problem
 - we can't do better than 'meh'

Pre-Classification

```
[1111000010100100001000000100000000000000000111111100001000]
```

Day 1

Night 1

Day 2

Night 2

- split the dataset into m groups
- implement m hidden layers for the perception and do **much** better than 'meh'
- Unfortunately, with sigmoids this is not possible (Vanishing Gradient)

Solution 2: We Go Deeper

Rectified Linear Activation Function (ReLU)

A piecewise function with constant derivative (linear). It's just a line with a cutoff!

Overhaul sigmoidals and use a ReLU instead

- Solve vanishing gradient problem → more hidden layers!
- unlock more sophisticated 'deep' learning algorithms (CNN, A2C)

Solution 2: We Go Deeper

Pros

- more informative and reliable output
- I can still code this myself (with numpy)

Cons

- more hidden layers are more expensive to train

PyTorch

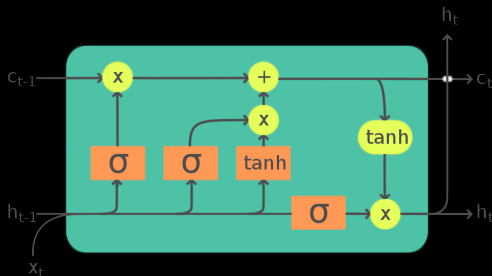
PyTorch Package for Python

“An open source machine learning framework that accelerates the path from research prototyping to production deployment”

PyTorch makes it easy to build things I don't understand!

- A2C - Actor Critic
 - value-based actor (Q-learning)
 - policy-based critic (perceptron)
- CNN - Convolution Neural Net
 - multilayered sparsely connected perceptron
 - GoTo choice for image classification
- LSTM - Long-Short Term Memory
 - recursive neural net with variable-time weight updates

Solution 3: LSTM



- Recursive (not feedforward) Neural Net
 - feedback connections allow for processing entire sequences of data - Wikipedia
- Sigmoidal - but somehow still works
 - solves Vanishing Gradient problem by allowing gradients to flow *unchanged* - Wikipedia

Solution 3: LSTM

Pros

- invented in 1995
- Bill Gates likes it
- OpenAI and Deepmind use it
- made for problems like ours

Cons

- might be expensive
- going to require more research
- I definitely couldn't make this without PyTorch